

Enhanced Base-Delta Compression with Memory Pooling

Overview

- **Base-Delta Compression** [Pekhimenko et. al., PACT'12] proposes a promising technique for increasing on chip cache capacity using **compression**.
- **B+ Δ** offers good compression but incurs an **additional access latency**.
- **B+ Δ** suffers **poor compressibility** when adjacent data in memory have **large value ranges**.
- **Observation**: Traditional compilers and memory-allocators are unaware of **B+ Δ** cache compression in hardware.
- **Key Idea**: Arrange data in memory to optimize B+ Δ compressibility.
- **Solution**: Recent literature on **Memory Pooling** and **Data Splitting** [Curial et. al., ISMM'08] and related work seem promising.

Motivation

Problem: Can we mitigate low compressibility cases for **B+ Δ** compression?

- Increase viability for B+ Δ implementation in hardware, and justify the extra access latency.
- Proposals like Memory Pooling and Data Splitting **already improve locality and reduce value range** in adjacent data values.
- **But they have not yet been applied to B+ Δ !**

Mechanisms

Basic Splitting-Pooling Example (64-bit)

Simple struct (a **node** perhaps)



In memory layout (**high range in adjacent values**)



After split-pool allocation (**much lower range**)



Proof of Concept Methodology

- To test the affect of splitting and pooling on **B+ Δ** compression, we manually restructured programs for optimal data layout.
- Ideally these pointer transformations will be implemented in the compiler.

Effect of Pooling on B+ Δ Compression Block Types

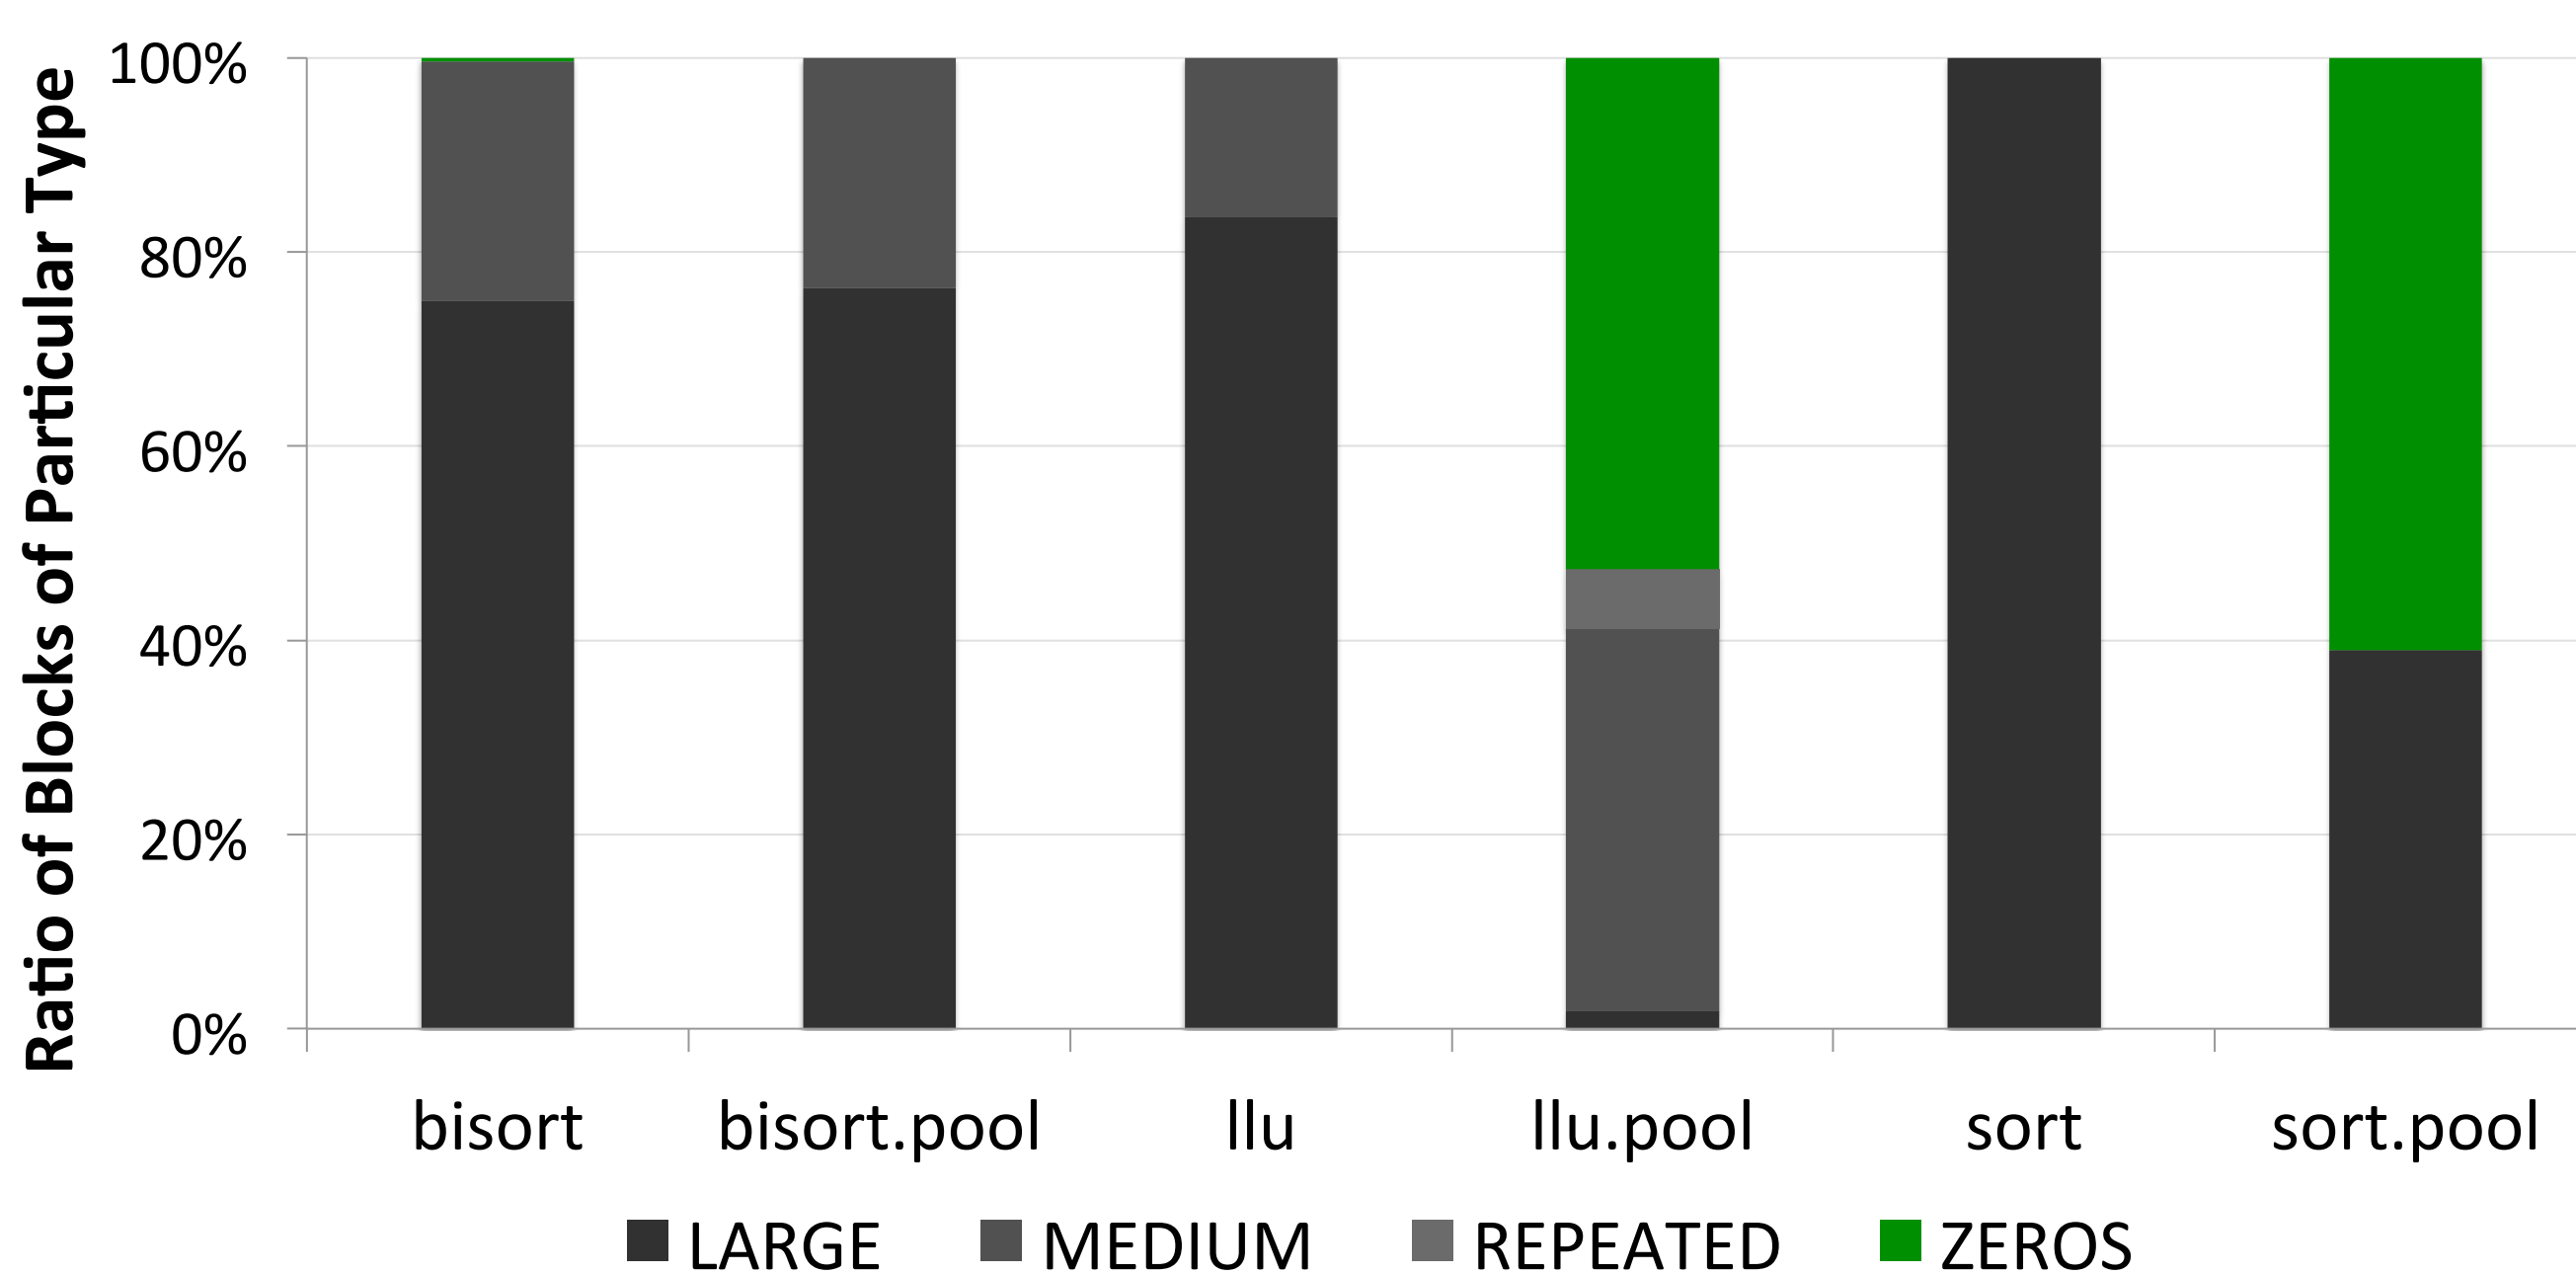
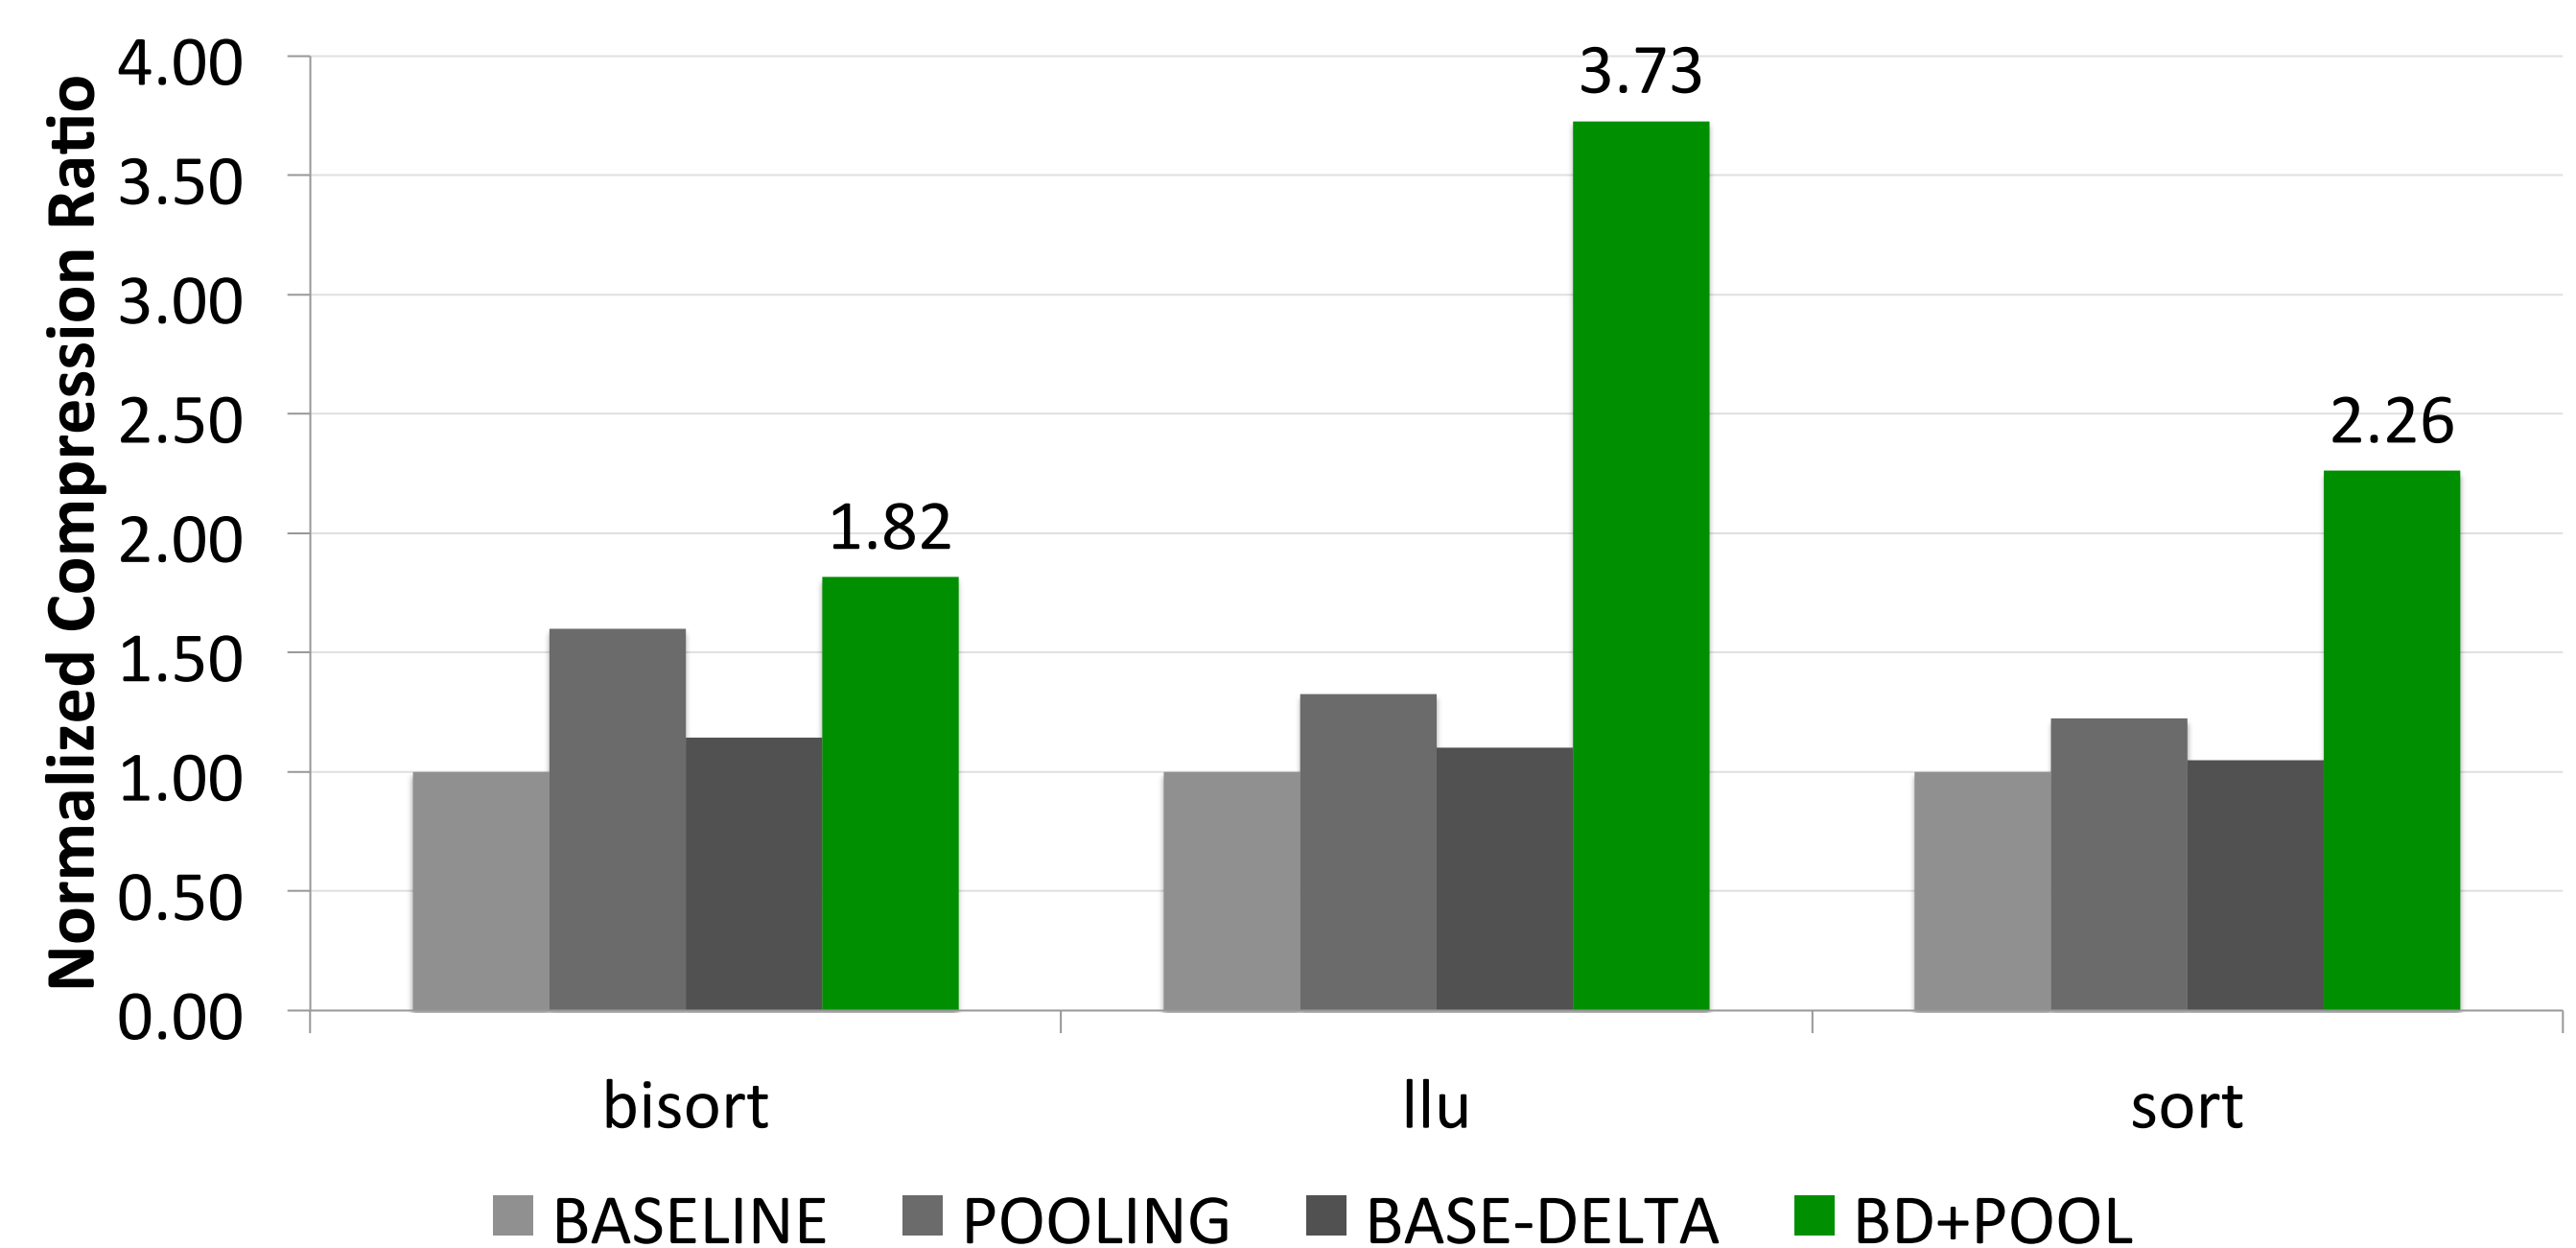


Figure 1. Each column shows the ratio of block-types for B+D compression with and without splitting and pooling. **Notice the large increase in 1-byte all-zero blocks, and general decrease of large, uncompressed blocks.**

B+ Δ compression on a 2MB, 16-way, 32BiB cache.

Results

Normalized Compression Ratio for Benchmarks



Conclusions

Min-Eviction: a novel replacement policy for the compressed cache

- Outperforms current state-of-the-art replacement policies
- First to consider both compressed block size and probability of reuse
- Simple to implement

Further Work:

- **Global Min-Eviction**: a global replacement policy for the compressed decoupled variable way cache that applies similar insight as Min-Eviction
- **Fairness** in compressed cache replacement
- **Multi-core evaluation and analysis** (see paper): 4% increase in normalized weighted speedup over LRU in heterogeneous workloads