

Enhanced Base-Delta Compression with Memory Pooling

Overview

- **Base-Delta Compression** [Pekhimenko et. al., PACT'12] proposes a promising technique for increasing on chip cache capacity using **compression**.
- **B+Δ** offers good compression but incurs an **additional access latency**.
- **B+Δ** suffers **poor compressibility** when adjacent data in memory have **large value ranges**.
- **Observation**: Traditional compilers and memory-allocators are unaware of **B+Δ** cache compression in hardware.
- **Key Idea**: Arrange data in memory to optimize B+Δ compressibility.
- **Solution**: Recent literature on [Memory Pooling, Data Splitting](#) [Curial et. al., ISMM'08] and related work seems promising.

Mechanisms

Basic Splitting-Pooling Example (64-bit)

Simple struct (a [node](#) perhaps)

FLAG (1B) VAL (4B) POINTER (8B)

In memory layout (**high range in adjacent values**)

... 0 100 0x77..0 1 108 0x77..4 0 93 0x77..C ...

After split-pool allocation (**much lower range**)

... 0 1 0 ... 0x77..0 0x77..4 0x77..C ... 100 108 93 ...

After **B+Δ** compression (**huge space savings**)

... B0 +0 +1 +0 ... B0x77..0 +0x0 +0x4 +0xC ... B100 +0 +8 -7 ...

Proof of Concept Methodology

- To test the affect of splitting and pooling on **B+Δ** compression, we manually restructured programs for optimal data layout. ([Later: implement pointer transformations in compiler](#))
- For this project, we focused on pointer based algorithms for benchmarks (**bisort** and **llu** – an aprx. for *Health*)

Effect of Pooling on B+Δ Compression Block Types

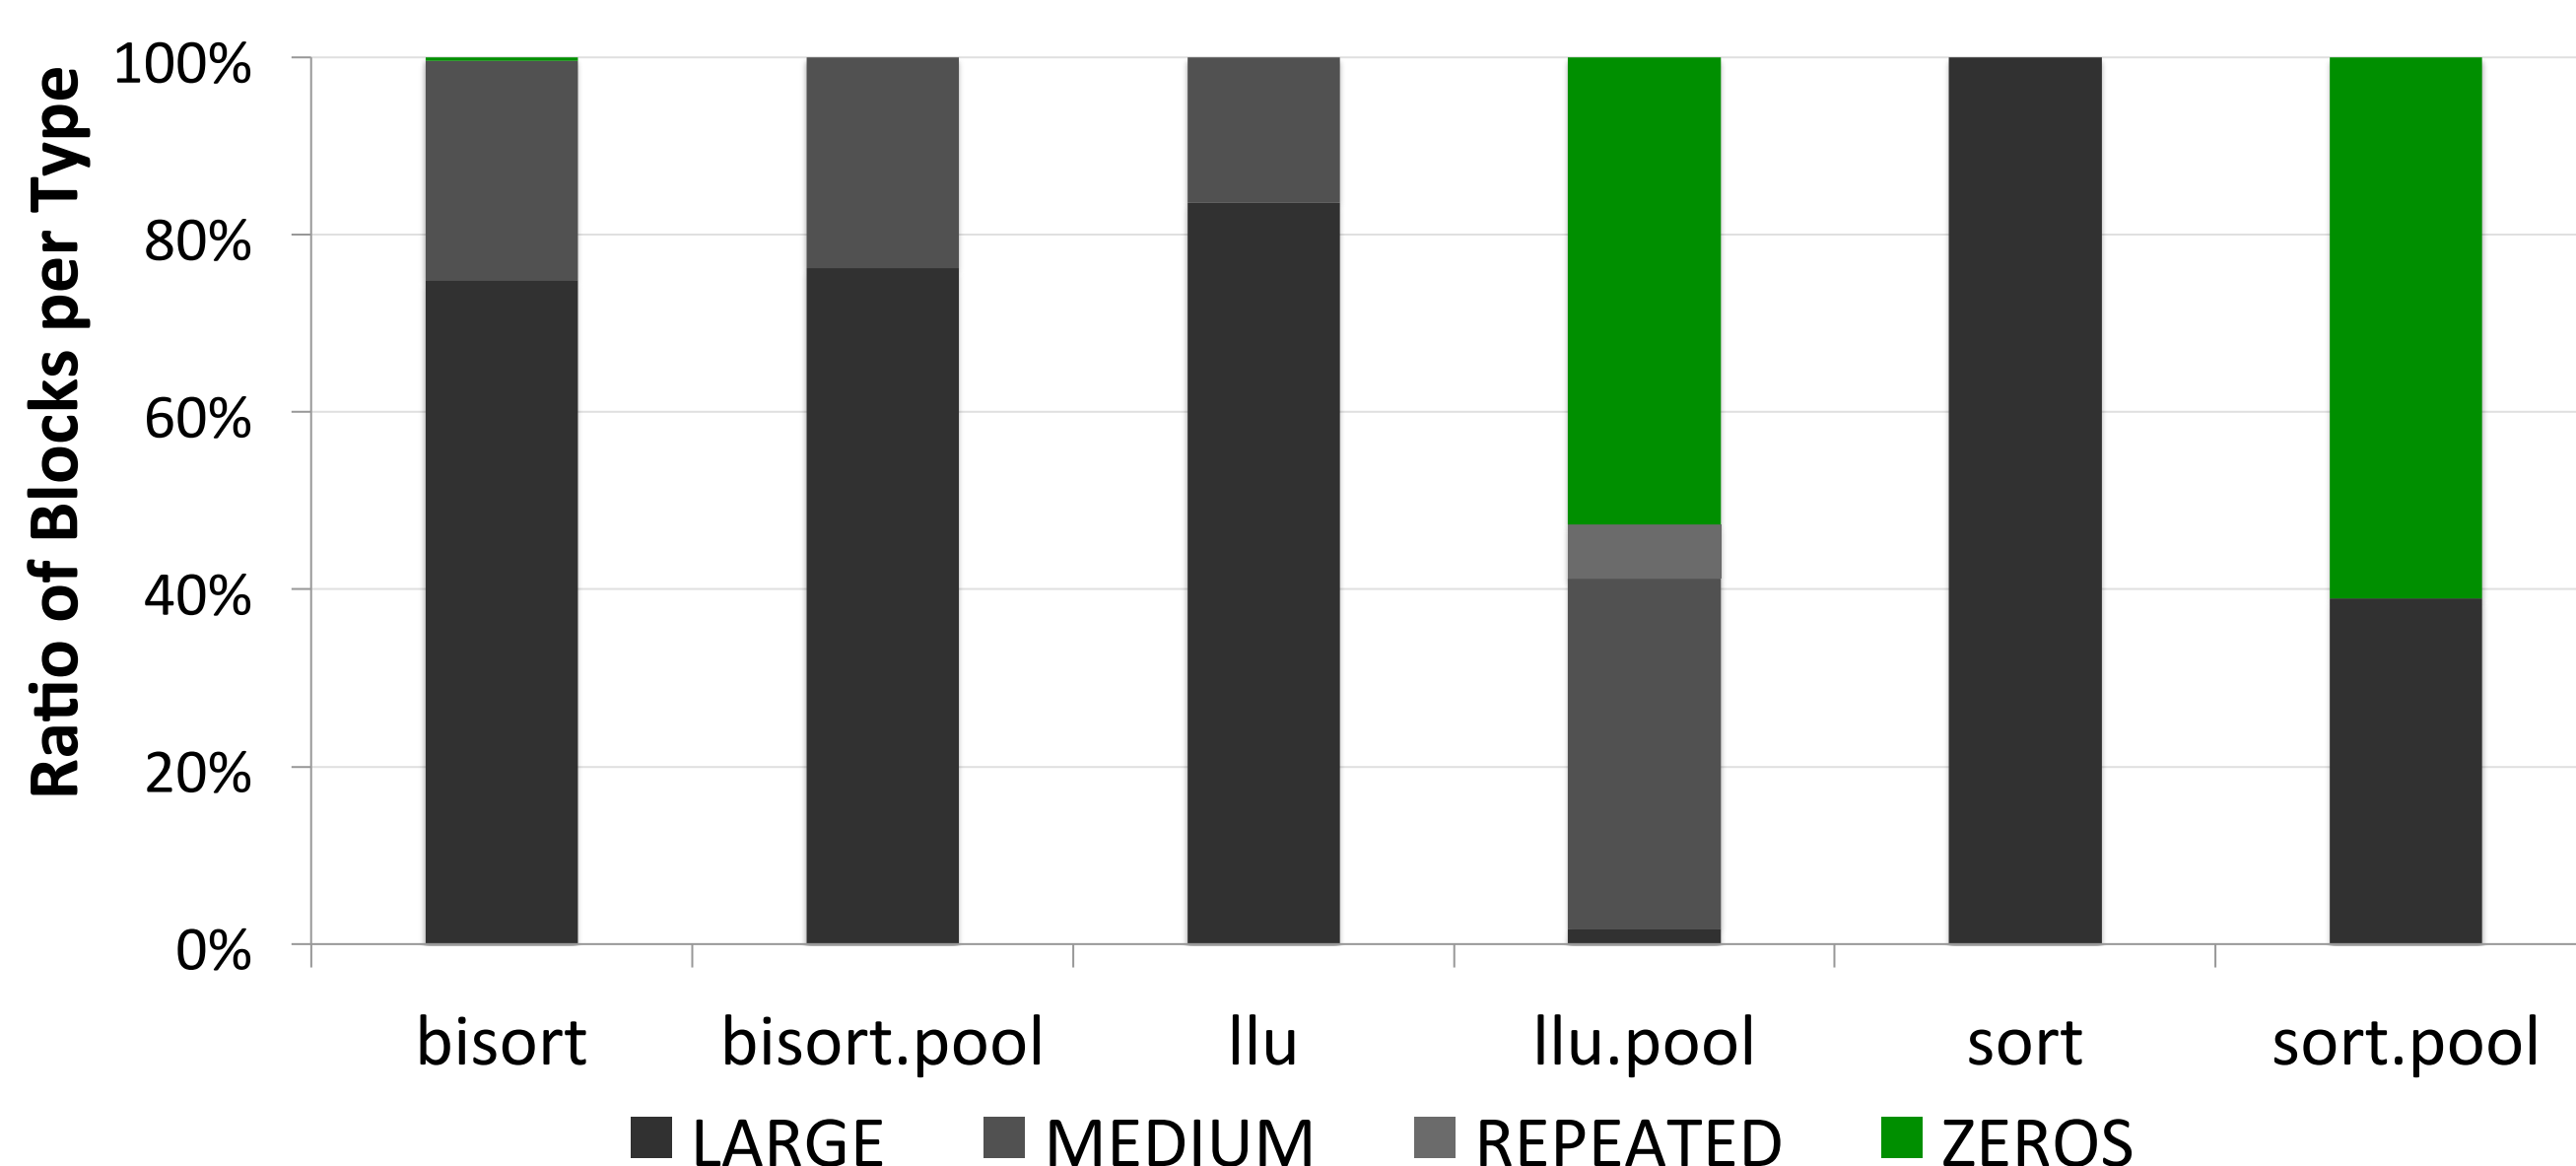


Figure 1. Each column shows the ratio of block-types for B+D compression with and without splitting and pooling. Notice the large increase in 1-byte all-zero blocks, and general decrease of large, uncompressed blocks. (1 MB, 16-way, 32-BiB BΔ-Cache)

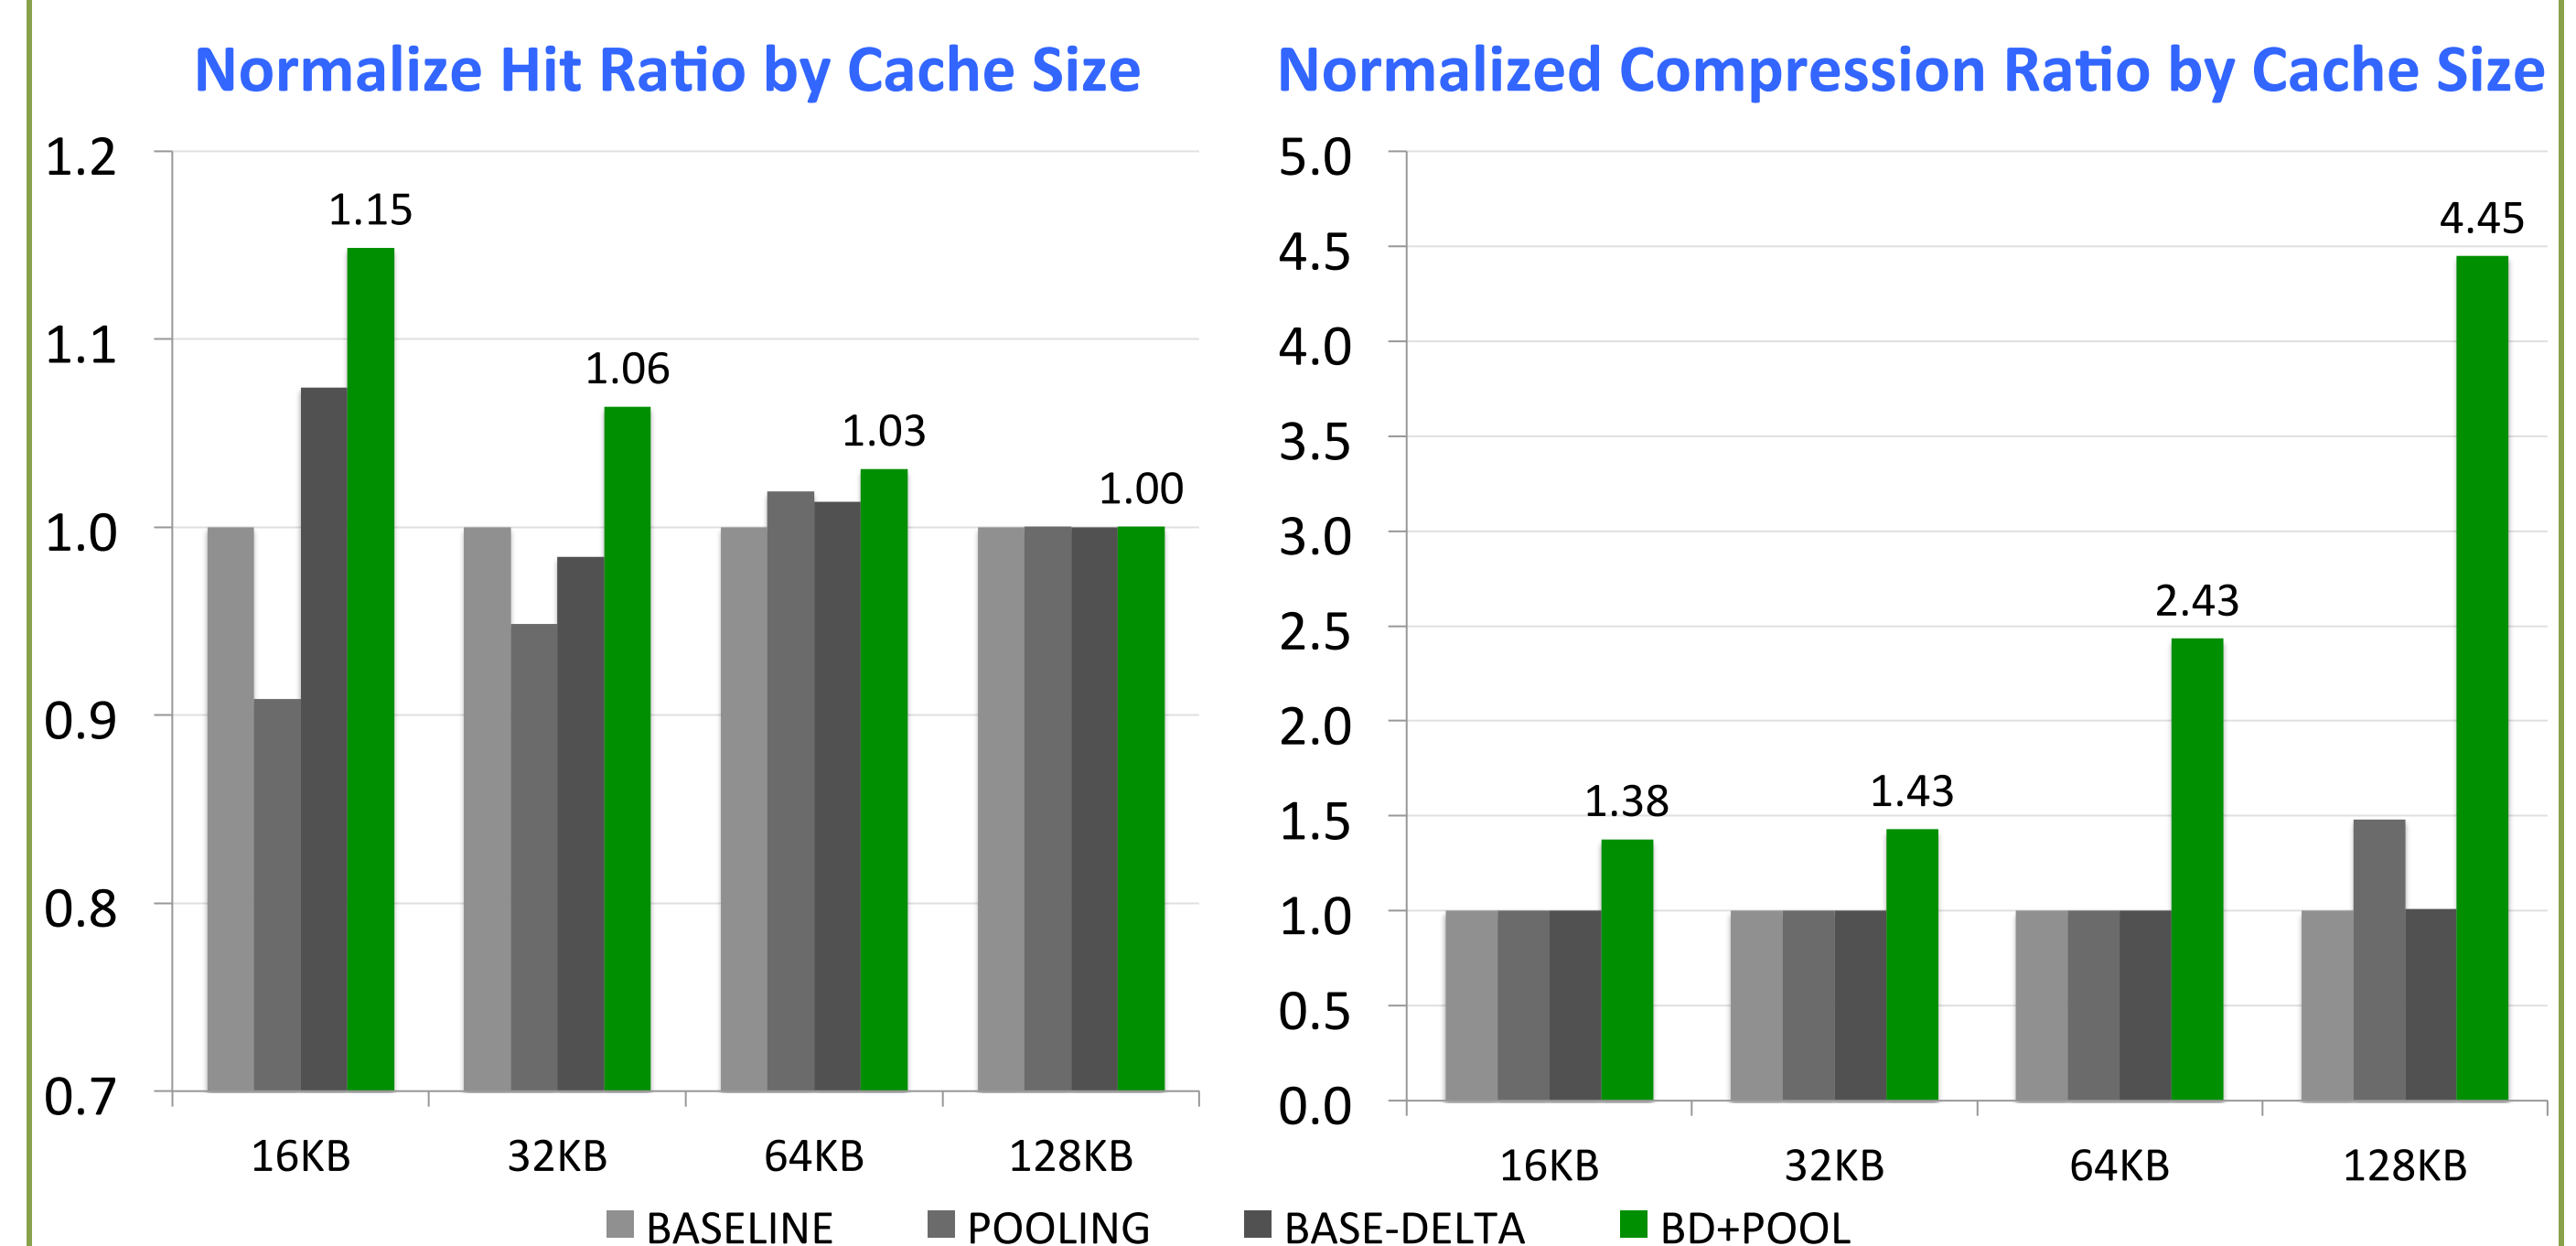
Motivation

Problem: Can we mitigate low compressibility cases for **B+Δ** compression?

- Increase viability for B+Δ implementation in hardware, and justify the extra access latency.
- Proposals like Memory Pooling and Data Splitting **already improve locality and reduce value range** in adjacent data values.
- But they have not yet been applied to **B+Δ**!

Results

Results for LLU micro-benchmark (working set ~117kb)



- **Improvement** in hit% from **fewer evictions** (more space)
- **B+Δ** alone reaches cache capacity for sizes < working set.
- **BD+POOL** still comes up with space savings!
- **Compression Ratio: 2.6x avg** (over LLU, TreeSort, ArraySort) **1.93x over just Split-Pool, 2.47x over only B+Δ**
- **Hit Rate: 8% avg. increase over micro-benchmarks.**
- Pointer based algorithms had poor locality.
- Expect multithreaded apps benefit from compression too.

Conclusions

BΔ-POOL: Strong improvement over baseline, pooling and base-delta

- Just proof of concept*
- Makes single base version of **BΔ** more viable.

*Recall that splitting and pooling was done by hand. [Safely](#) splitting-pooling in the compiler is not always possible.

Further Work

- Implement pointer transformations in LLVM.
 - Run benchmarks on cycle accurate BΔ-simulator.
 - Incorporate data from standard benchmarks.
 - Multithreaded environments.
 - Interaction with non-traditional LRU policies.
- See CARP by Huberty et. al.