

Enhanced Base-Delta Compression with Memory Pooling

Overview

- Base-Delta Compression** [Pekhimenko et. al., PACT'12] proposes a promising technique for increasing on chip cache capacity using **compression**.
- B+Δ** offers good compression but incurs an **additional access latency**.
- B+Δ** suffers **poor compressibility** when adjacent data in memory have **large value ranges**.
- Observation**: Traditional compilers and memory-allocators are unaware of **B+Δ** cache compression in hardware.
- Key Idea**: Arrange data in memory to optimize B+Δ compressibility.
- Solution**: Recent literature on **Memory Pooling, Data Splitting** [Curial et. al., ISMM'08] and related work seems promising.

Mechanisms

Basic Splitting-Pooling Example (64-bit)

Simple struct (a **node** perhaps)



In memory layout (**high range in adjacent values**)



After split-pool allocation (**much lower range**)



After **B+Δ** compression (**huge space savings**)



Proof of Concept Methodology

- To test the affect of splitting and pooling on **B+Δ** compression, we manually restructured programs for optimal data layout. (**Later: implement pointer transformations in compiler**)
- For this project, we focused on pointer based algorithms for benchmarks (**bisort** and **llu** – an aprx. for *Health*)

Effect of Pooling on B+Δ Compression Block Types

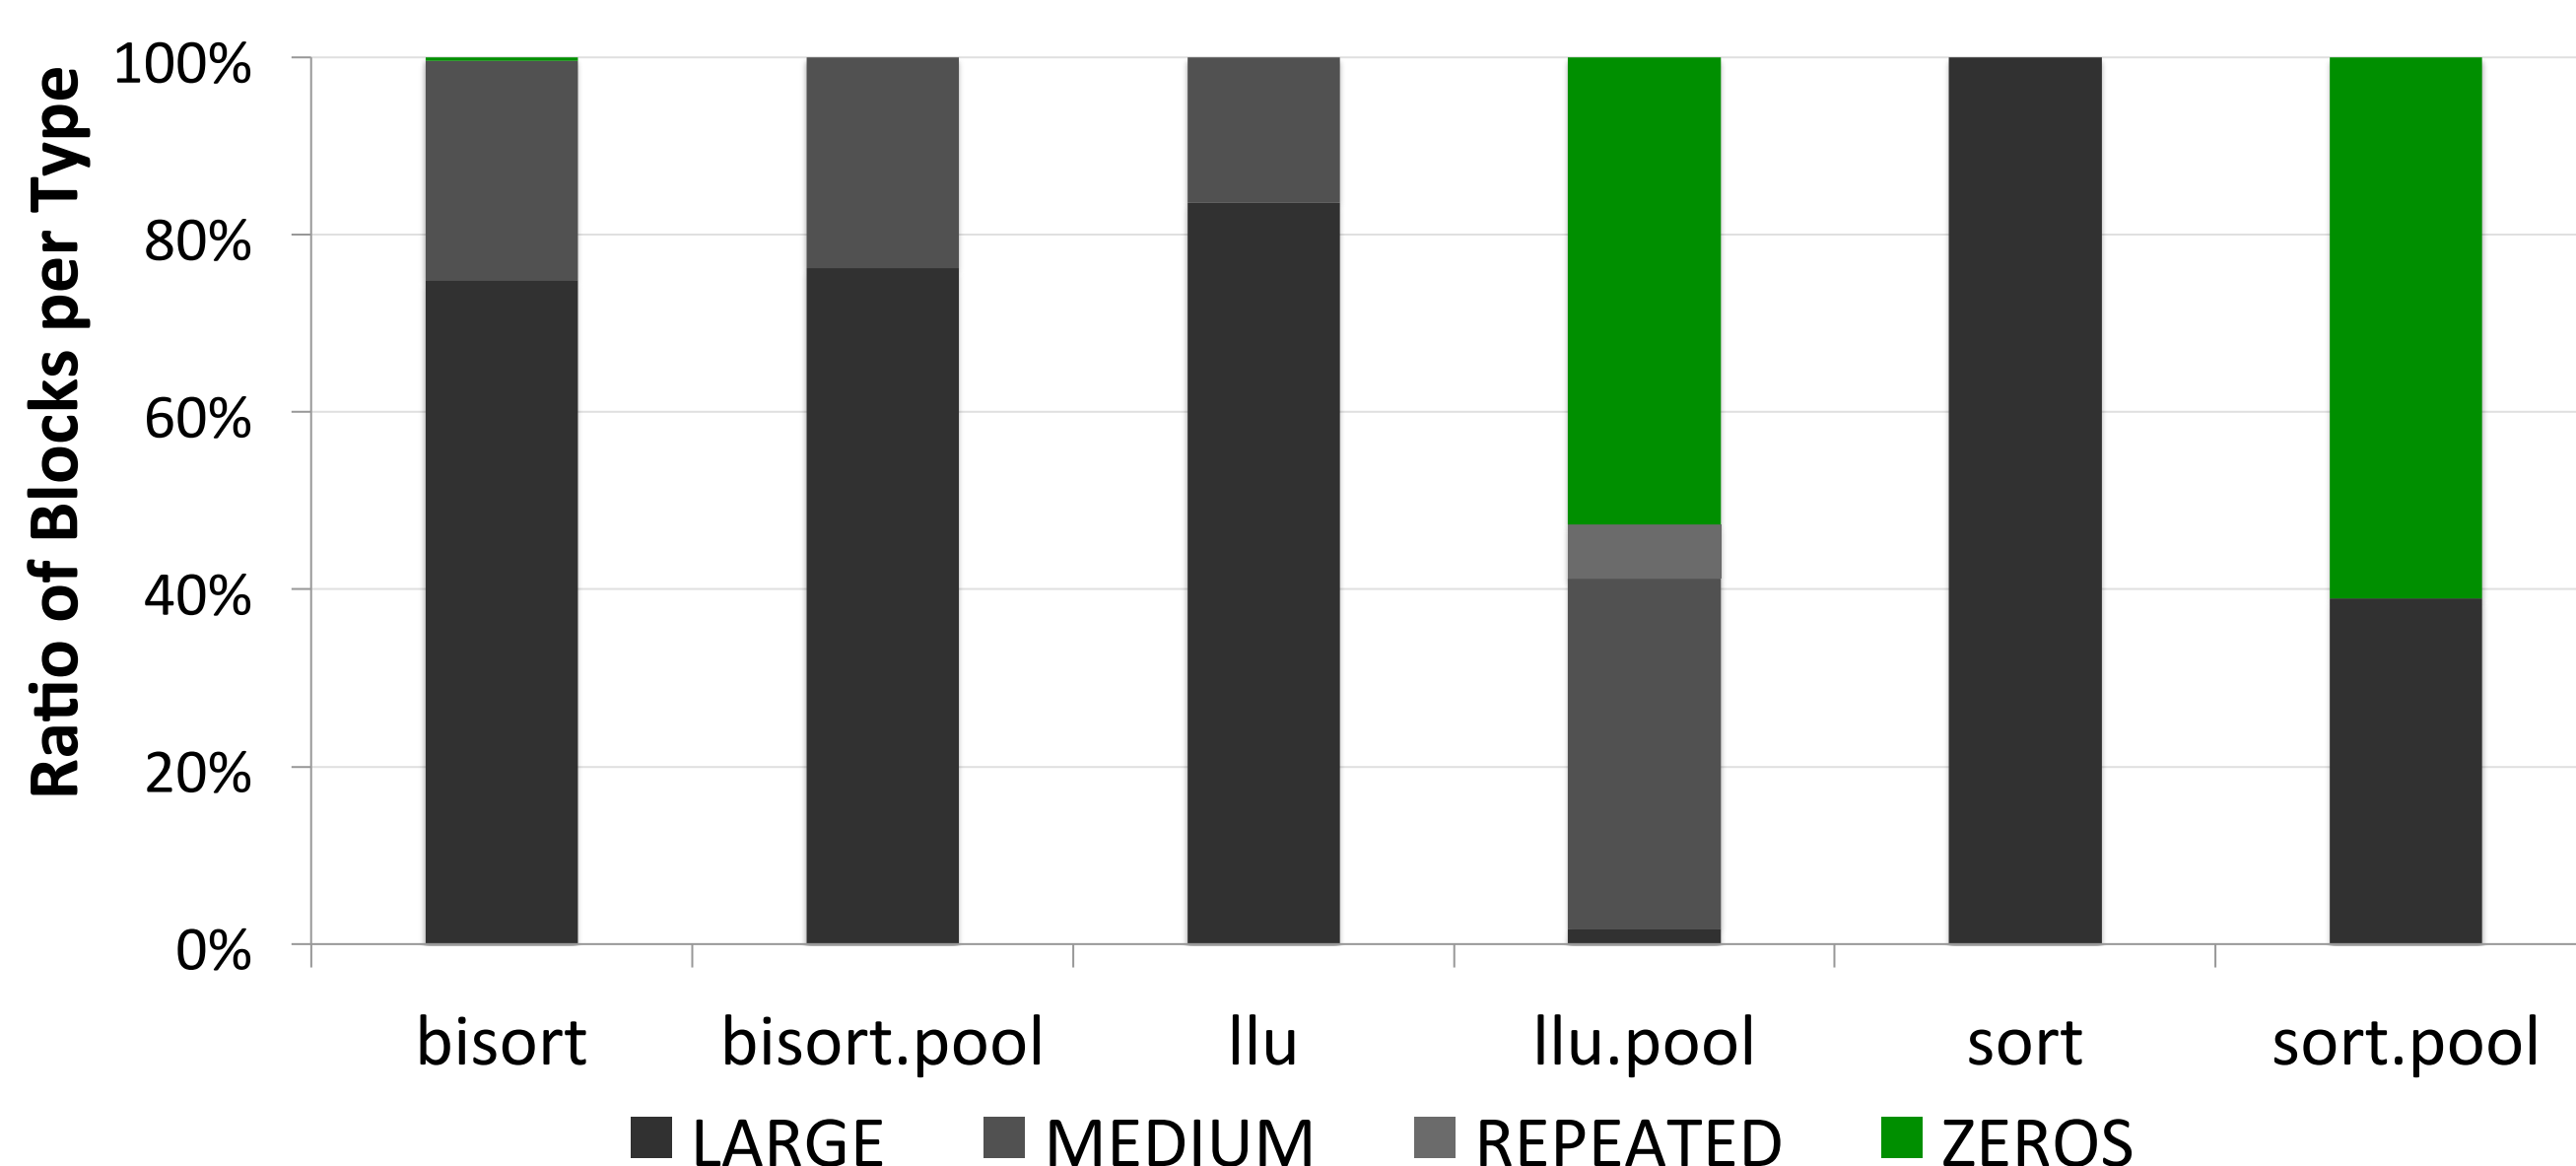


Figure 1. Each column shows the ratio of block-types for B+D compression with and without splitting and pooling. **Notice the large increase in 1-byte all-zero blocks, and general decrease of large, uncompressed blocks.** (1 MB, 16-way, 32-BiB BΔ-Cache)

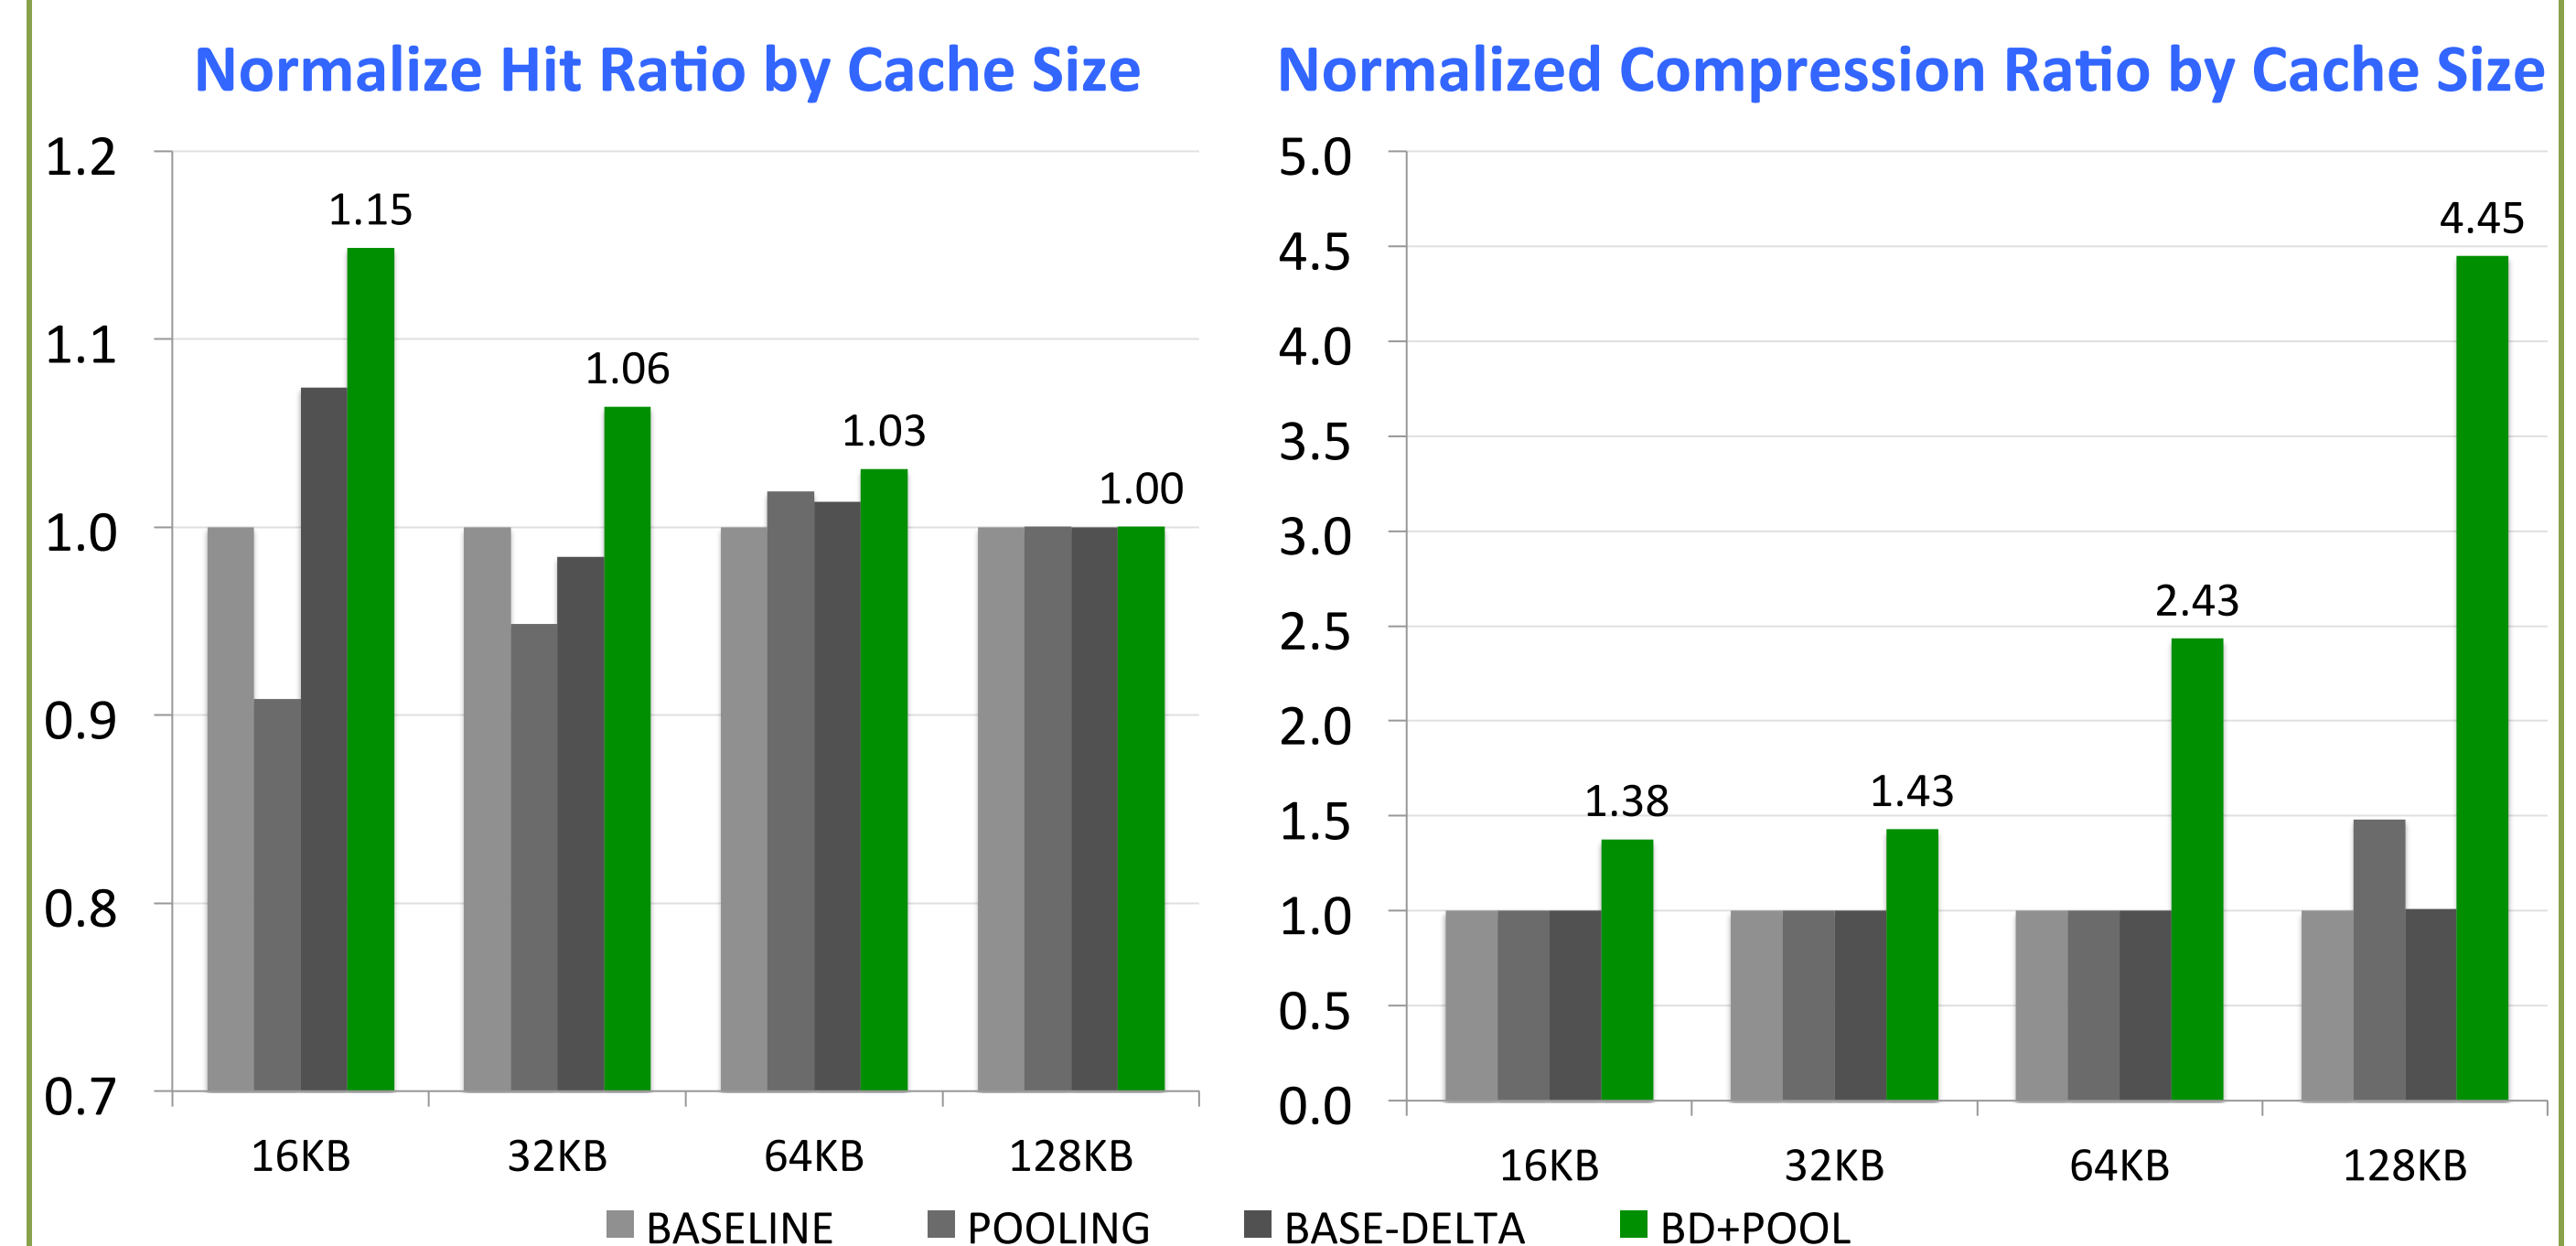
Motivation

Problem: Can we mitigate low compressibility cases for **B+Δ** compression?

- Increase viability for B+Δ implementation in hardware, and justify the extra access latency.
- Proposals like Memory Pooling and Data Splitting **already improve locality and reduce value range** in adjacent data values.
- But they have not yet been applied to B+Δ!**

Results

Results for LLU micro-benchmark (working set ~117kb)



- Improvement** in hit% from **fewer evictions** (more space)
- B+Δ** alone reaches cache capacity for sizes < working set.
- BD+POOL** still comes up with space savings!
- Compression Ratio: 2.6x avg** (over LLU, TreeSort, ArraySort) **1.93x over just Split-Pool, 2.47x over only B+Δ**
- Hit Rate: 8% avg. increase over micro-benchmarks.**
- Pointer based algorithms had poor locality.
- Expect multithreaded apps benefit from compression too.

Conclusions

BΔ-POOL: Strong improvement over baseline, pooling and base-delta

- Just proof of concept***
- Makes single base version of BΔI more viable.**

*Recall that splitting and pooling was done by hand. **Safely** splitting-pooling in the compiler is not always possible.

Further Work

- Implement pointer transformations in LLVM.
 - Run benchmarks on cycle accurate BΔ-simulator.
 - Incorporate data from standard benchmarks.
 - Multithreaded environments.
 - Interaction with non-traditional LRU policies.
- See CARP by Huberty et. al.